

به نام خداوند بخشنده و مهربان



دانشگاه صنعتی امیرکبیر  
( پلی تکنیک تهران )

## درس یادگیری ماشین

گزارش تمرین شش

استاد درس: دکتر سعید شیری قیداری

نام دانشجو:

روزبه قاسمی ۹۵۳۱۴۲۴

زمستان ۱۳۹۷

## سوال ۱

طبق مطالبی که در سایتی که معرفی کرده بودید یعنی سایت [deeplearning.ir](http://deeplearning.ir) مطلب رو به رو [کلینک کنید!](#) را پیدا کردم که توضیح داده بود می‌توانیم به دیتابیس های `hdf5` یا `lmbd` یا `leveldb` تبدیل کنیم.

## سوال ۲

در ابتدا کد را از منبع آن یعنی سایتی که مشخص کرده بودید دانلود کردم که کد آن به شکل زیر است:

```
from urllib.request import urlretrieve
from os.path import isfile, isdir
from tqdm import tqdm
import tarfile

cifar10_dataset_folder_path = 'cifar-10-batches-py'

class DLProgress(tqdm):
    last_block = 0

    def hook(self, block_num=1, block_size=1, total_size=None):
        self.total = total_size
        self.update((block_num - self.last_block) * block_size)
        self.last_block = block_num

    if not isfile('cifar-10-python.tar.gz'):
        with DLProgress(unit='B', unit_scale=True, miniters=1, desc='CIFAR-10
Dataset') as pbar:
            urlretrieve(
                'https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz',
                'cifar-10-python.tar.gz',
                pbar.hook)

    if not isdir(cifar10_dataset_folder_path):
        with tarfile.open('cifar-10-python.tar.gz') as tar:
            tar.extractall()
            tar.close()
```

سپس در مرحله بعد عکس های موجود در این مجموعه داده را برچسب گذاری می‌کنیم و مشخص می‌کنیم از هر برچسب چند تا عکس داریم و از هر کدام یه نمونه خروجی می‌گیریم که کد آن به شکل زیر است:

```
import pickle
import matplotlib.pyplot as plt

LABEL_NAMES = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog',
                'horse', 'ship', 'truck']

def load_cifar10_batch(batch_id):
    """
    Load a batch of the dataset
```

```

"""
with open(cifar10_dataset_folder_path + '/data_batch_'
+ str(batch_id), mode='rb') as file:
    batch = pickle.load(file, encoding='latin1')

    features = batch['data'].reshape((len(batch['data']), 3, 32, 32)).transpose(0, 2, 3,
    1)
    labels = batch['labels']

    return features, labels

def display_stats(features, labels, sample_id):
    """
    Display Stats of the the dataset
    """

    if not (0 <= sample_id < len(features)):
        print('{} samples in batch {}. {} is out of range.'
        .format(len(features), batch_id, sample_id))
        return None

    print('\nStats of batch {}'.format(batch_id))
    print('Samples: {}'.format(len(features)))
    print('Label Counts: {}'.format(dict(zip(*np.unique(labels, return_counts=True))))
    print('First 20 Labels: {}'.format(labels[:20]))

    sample_image = features[sample_id]
    sample_label = labels[sample_id]

    print('\nExample of Image {}'.format(sample_id))
    print('Image - Min Value: {} Max Value: {}'.format(sample_image.min(),
    sample_image.max()))
    print('Image - Shape: {}'.format(sample_image.shape))
    print('Label - Label Id: {} Name: {}'.format(sample_label, LABEL_NAMES[sample_label]))
    plt.axis('off')
    plt.imshow(sample_image)
    plt.show()

    %matplotlib inline
    %config InlineBackend.figure_format = 'retina'

import numpy as np

for batch_id in range(1,6):
    features, labels = load_cifar10_batch(batch_id)
    for image_id in range(0,2):
        display_stats(features, labels, image_id)

del features, labels # free memory

```

که خروجی های آن مانند زیر است:

```

Stats of batch 1:
Samples: 10000
Label Counts: {0: 1005, 1: 974, 2: 1032, 3: 1016, 4: 999, 5: 937, 6: 1030, 7:
1001, 8: 1025, 9: 981}
First 20 Labels: [6, 9, 9, 4, 1, 1, 2, 7, 8, 3, 4, 7, 7, 2, 9, 9, 9, 3, 2, 6]

```

Example of Image 0:  
Image - Min Value: 0 Max Value: 255  
Image - Shape: (32, 32, 3)  
Label - Label Id: 6 Name: frog



حال در ادامه می بینیم که پیاده سازی من سه لایه است اما ابتدا به قسمت پیش پردازش خواسته شده می پردازیم:

(A)

حال عملیات هایی که برای پیش پردازش داده ها انجام دادم به شکل زیر است:

۱- نرمالایز کردن داده ها

```
def normalize(x):  
    # Each pixel has three channels - Red, Green and Blue.  
    # Each channel is an int between 0 and 255 (8-bit color scheme).  
    return np.array(x) / 255.0
```

۲- One-hot-encoding

با توجه به نکته ای در این سایت [one-hot-encoding](#) گفته شده است اینکار را کردم:

```
def one_hot_encode(x):  
    one_hot_encoded = np.zeros((len(x), 10))
```

```
for i in range(len(x)):
    one_hot_encoded[i][x[i]] = 1.0
return one_hot_encoded
```

## ۳- mirroring

از دیگر پیش پردازش ها این عمل بود:

```
from PIL import Image

def flip_image(image_path, saved_location):
    """
    Flip or mirror the image
    """
    @param image_path: The path to the image to edit
    @param saved_location: Path to save the cropped image
    """
    image_obj = Image.open(image_path)
    rotated_image = image_obj.transpose(Image.FLIP_LEFT_RIGHT)
    rotated_image.save(saved_location)
    rotated_image.show()

    if name == 'main':
        image = 'mantis.png'
        flip_image(image, 'flipped_mantis.jpg')
```

**(B)**

در این قسمت باید بگویم اگر مقدار **iteration** ها را زیاد یا کم کنیم چه اتفاقی می افتد. طبق مطالبی که در سایت های مختلف و یوتیوب دیدم این بود که اگر از یک مقداری بیش از حد شود باعث می شود **Overfitting** رخ دهد و اگر از یه حدی کم شود باعث می شود **underfitting** رخ دهد در نتیجه باید یک مقدر مشخصی باشد.

همچنین اگر مقدار **batch size** باید متناسب با مقدار **RAM** باشد و اگر بتواند **RAM** پشتیبانی کند اگر زیاد کنیم باعث می شود در هر مرحله مقدار بیشتری محاسبات می کند و به طبع سرعت بالاتری نیز دارد.

حال شبکه عصبی که طراحی کردم را ایجاد می کنم:

ورودی آن به شکل زیر است:

```
import tensorflow as tf

def neural_net_image_input(image_shape):
    return tf.placeholder(tf.float32, shape=(None, image_shape[0], image_shape[1],
image_shape[2]), name='x')
```

```
def neural_net_label_input(n_classes):
    return tf.placeholder(tf.float32, shape=(None, n_classes), name='y')

def neural_net_keep_prob_input():
    return tf.placeholder(tf.float32, name='keep_prob')
```

حال لایه convolutional و max pooling آن را ایجاد کردیم:

```
def conv2d_maxpool(x_tensor, conv_num_outputs, conv_ksize, conv_strides, pool_ksize,
pool_strides):

    # Create filter dimensions
    filter_height, filter_width, in_channels, out_channels = \
    conv_ksize[0], conv_ksize[1], x_tensor.get_shape().as_list()[3], conv_num_outputs
    conv_filter = [filter_height, filter_width, in_channels, out_channels]

    # Create weights and bias
    weights = tf.Variable(tf.truncated_normal(conv_filter, stddev=0.05))
    bias = tf.Variable(tf.truncated_normal([conv_num_outputs], stddev=0.05))

    # Create strides
    strides=(1,conv_strides[0], conv_strides[1], 1)

    # Bind all together to create the layer
    conv = tf.nn.conv2d(x_tensor, weights, strides, padding='SAME')
    conv = tf.nn.bias_add(conv, bias)

    # Create ksize
    ksize = (1, pool_ksize[0], pool_ksize[1], 1)

    # Create strides
    strides=(1,pool_strides[0], pool_strides[1], 1)

    pool = tf.nn.max_pool(conv, ksize, strides, padding='SAME')

    print('Convolutional layer with conv_num_outputs:',conv_num_outputs,
    'conv_ksize:', conv_ksize,
    'conv_strides:', conv_strides,
    'pool_ksize:',pool_ksize,
    'pool_strides', pool_strides)
    print('layer input shape', x_tensor.get_shape().as_list(),
    'layer output shape', pool.get_shape().as_list())

    return pool
```

حال لایه flatten و fully connected آن را ایجاد می‌کنیم:

```
def flatten(x_tensor):
    height, width, channels = x_tensor.get_shape().as_list()
```

```

net = tf.reshape(x_tensor, shape=[-1, height * width * channels])
print('flatten shape', net.get_shape().as_list())
return net

def fully_conn(x_tensor, num_outputs):
    _, size = x_tensor.get_shape().as_list()
    weights = tf.Variable(tf.truncated_normal([size, num_outputs], stddev=0.05))
    bias = tf.Variable(tf.truncated_normal([num_outputs], stddev=0.05))

    fully_connected = tf.add(tf.matmul(x_tensor, weights), bias)

    print('layer input shape', x_tensor.get_shape().as_list(),
          'layer output shape', fully_connected.get_shape().as_list())

    return fully_connected

```

حالا معماری شبکه عصبی را مشخص می‌کنیم:

```

def conv_net(input_x, keep_probability):

    net = conv2d_maxpool(input_x, 32, (7,7), (2,2), (2,2), (2,2))
    net = conv2d_maxpool(net, 64, (3,3), (1,1), (2,2), (2,2))
    net = conv2d_maxpool(net, 128, (2,2), (1,1), (2,2), (2,2))

    net = flatten(net)
    net = tf.nn.dropout(net, keep_probability)
    net = fully_conn(net, 1024)
    net = tf.nn.dropout(net, keep_probability)
    net = fully_conn(net, 128)
    net = fully_conn(net, 10)

    return net

```

حالا شبکه عصبی را می‌سازیم:

```

#####
## Build the Neural Network ##
#####

# Remove previous weights, bias, inputs, etc..
tf.reset_default_graph()

```

```

IMAGE_SHAPE = (32, 32, 3)
LABELS_COUNT = 10

# Inputs
x = neural_net_image_input(IMAGE_SHAPE)
y = neural_net_label_input(LABELS_COUNT)
keep_prob = neural_net_keep_prob_input()

# Model
logits = conv_net(x, keep_prob)

# Name logits Tensor, so that it can be loaded from disk after training
logits = tf.identity(logits, name='logits')

# Loss and Optimizer
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=logits,
labels=y))
optimizer = tf.train.AdamOptimizer(learning_rate=0.001).minimize(cost)

# Accuracy
correct_pred = tf.equal(tf.argmax(logits, 1), tf.argmax(y, 1))
accuracy = tf.reduce_mean(tf.cast(correct_pred, tf.float32), name='accuracy')

```

خروجی:

```

Convolutional layer with conv_num_outputs: 32 conv_ksize: (7, 7)
conv_strides: (2, 2) pool_ksize: (2, 2) pool_strides (2, 2)
layer input shape [None, 32, 32, 3] layer output shape [None, 8, 8, 32]
Convolutional layer with conv_num_outputs: 64 conv_ksize: (3, 3)
conv_strides: (1, 1) pool_ksize: (2, 2) pool_strides (2, 2)
layer input shape [None, 8, 8, 32] layer output shape [None, 4, 4, 64]
Convolutional layer with conv_num_outputs: 128 conv_ksize: (2, 2)
conv_strides: (1, 1) pool_ksize: (2, 2) pool_strides (2, 2)
layer input shape [None, 4, 4, 64] layer output shape [None, 2, 2, 128]
flatten shape [None, 512]
layer input shape [None, 512] layer output shape [None, 1024]
layer input shape [None, 1024] layer output shape [None, 128]
layer input shape [None, 128] layer output shape [None, 10]

```

سپس نتایج آنرا نشان می‌دهم:

```

valid_features, valid_labels = pickle.load(open('preprocess_validation.p', mode='rb'))

def print_stats(session, feature_batch, label_batch, cost, accuracy):
    batch_loss = session.run(cost, feed_dict=\
    {x:feature_batch, y:label_batch, keep_prob:1.0})
    batch_accuracy = session.run(accuracy, feed_dict=\

```



```
{x:valid_features, y:valid_labels, keep_prob:1.0})

print('batch loss is : ', batch_loss)
print('batch accuracy is : ', batch_accuracy)
```

حال من Hyperparameter های آنرا مشخص کردم:

```
#For my laptop I selected:
epochs = 50
batch_size = 256

# keep probability of 70%
keep_probability = 0.5
```

```
#a couple of helper functions for loading a single batch
def batch_features_labels(features, labels, batch_size):
    """
    Split features and labels into batches
    """
    for start in range(0, len(features), batch_size):
        end = min(start + batch_size, len(features))
        yield features[start:end], labels[start:end]

def load_preprocess_training_batch(batch_id, batch_size):
    """
    Load the Preprocessed Training data and return them in batches of <batch_size> or
    less
    """
    filename = 'preprocess_batch_' + str(batch_id) + '.p'
    features, labels = pickle.load(open(filename, mode='rb'))

    # Return the training data in batches of size <batch_size> or less
    return batch_features_labels(features, labels, batch_size)
```

### سوال ۳

در این سوال از ما خواسته شده بررسی کنیم که اگر تعداد پارامترها ثابت باشد، عمق شبکه را اگر زیاد یا کم کنیم چه تاثیری می‌گذارد. پاسخ در این است که اگر بیش از حد عمق را زیاد کنیم هم باعث **overfitting** می‌شود و هم این که باعث می‌شود محاسبات بیش از حد انجام بدهیم که این یعنی از منابع که شامل سخت افزار سیستم می‌شود، بیهوده استفاده کرده‌ایم.

(A)

برای اینکار کافی است از سه لایه ای که من طراحی کنیم فقط یک لایه رو در نظر بگیریم و همین کد را دوباره اجرا کنیم.

(B)

برای این قسمت گفته اید حداقل یک لایه و خب من همان سه لایه خودم را در نظر می گیریم.

(C)

از نتایج بدست آمده نتیجه گرفتیم accuracy به طور محسوسی کاهش یافت!

## سوال ۴

در این سوال نیاز است لایه های مان را فریز کنیم:

دستورات آن به شکل زیر است:

```
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
from __future__ import print_function
import keras
from keras.preprocessing.image import ImageDataGenerator, load_img
```

حال برای اینکه بتوانیم دیتاست تصاویر را بخوانیم آن را در google drive خود قرار می دهیم:

```
# i have got this code from this link(way number 3): https://towardsdatascience.com/3-ways-to-load-csv-files-into-colab-7c14fcbdc92
# Code to read csv file into Colaboratory:

!pip install -U -q PyDrive
from pydrive.auth import GoogleAuth
from pydrive.drive import GoogleDrive
from google.colab import auth
from oauth2client.client import GoogleCredentials
```

خروجی:

```
100% |████████████████████████████████████████| 993kB 20.3MB/s
Building wheel for PyDrive (setup.py) ... done
```

```
#Authenticate and create the PyDrive client.
auth.authenticate_user()
gauth = GoogleAuth()
gauth.credentials = GoogleCredentials.get_application_default()
drive = GoogleDrive(gauth)

link = 'https://drive.google.com/open?id=1t1-PPc5mO6OvlilY6ffkFrOJoqKSk3RJ'
fluff, id = link.split('=')

print (id) # Verify that you have everything after '='
```

حال آدرس فولدر های train و validation را به آن می‌دهیم:

```
train_dir = ' https://drive.google.com/open?id=15COsuoR3zbfiToquwB49d5j-epI5K0I2'
validation_dir = '
https://drive.google.com/open?id=1kNo13cgcaDsvt_NC71WQIbDuYwhE6wbY'
image_size = 224
```

حال مدل مان را بدست می‌آوریم:

```
from keras.applications import VGG16

#Load the VGG model
vgg_conv = VGG16(weights='imagenet', include_top=False, input_shape=(image_size,
image_size, 3))

# Freeze all the layers
for layer in vgg_conv.layers[:]:
    layer.trainable = False

# Check the trainable status of the individual layers
for layer in vgg_conv.layers:
    print(layer, layer.trainable)

from keras import models
from keras import layers
from keras import optimizers

# Create the model
model = models.Sequential()

# Add the vgg convolutional base model
model.add(vgg_conv)

# Add new layers
model.add(layers.Flatten())
model.add(layers.Dense(1024, activation='relu'))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(3, activation='softmax'))

# Show a summary of the model. Check the number of trainable parameters
model.summary()
```

```

<keras.engine.topology.InputLayer object at 0x7f2eed2f01d0> False
<keras.layers.convolutional.Conv2D object at 0x7f2eed2f0240> False
<keras.layers.convolutional.Conv2D object at 0x7f2eed2f0390> False
<keras.layers.pooling.MaxPooling2D object at 0x7f2eed2f0518> False
<keras.layers.convolutional.Conv2D object at 0x7f2eed29fa58> False
<keras.layers.convolutional.Conv2D object at 0x7f2eed29f860> False
<keras.layers.pooling.MaxPooling2D object at 0x7f2eed23f3c8> False
<keras.layers.convolutional.Conv2D object at 0x7f2eed260710> False
<keras.layers.convolutional.Conv2D object at 0x7f2eed260048> False
<keras.layers.convolutional.Conv2D object at 0x7f2eed203f28> False
<keras.layers.pooling.MaxPooling2D object at 0x7f2eed225da0> False
<keras.layers.convolutional.Conv2D object at 0x7f2eed1c8278> False
<keras.layers.convolutional.Conv2D object at 0x7f2eed1c8898> False
<keras.layers.convolutional.Conv2D object at 0x7f2eed1e97f0> False
<keras.layers.pooling.MaxPooling2D object at 0x7f2eed17c240> False
<keras.layers.convolutional.Conv2D object at 0x7f2eed19d860> False
<keras.layers.convolutional.Conv2D object at 0x7f2eed19d128> False
<keras.layers.convolutional.Conv2D object at 0x7f2eed13d780> False
<keras.layers.pooling.MaxPooling2D object at 0x7f2eed15fda0> False

```

Layer (type)	Output Shape	Param #
vgg16 (Model)	(None, 7, 7, 512)	14714688
flatten_1 (Flatten)	(None, 25088)	0
dense_1 (Dense)	(None, 1024)	25691136
dropout_1 (Dropout)	(None, 1024)	0
dense_2 (Dense)	(None, 3)	3075
Total params: 40,408,899		
Trainable params: 25,694,211		
Non-trainable params: 14,714,688		

حالا باید مدل را آموزش دهیم:

```

# No Data augmentation
train_datagen = ImageDataGenerator(rescale=1./255)
validation_datagen = ImageDataGenerator(rescale=1./255)

# Change the batchsize according to your system RAM
train_batchsize = 100
val_batchsize = 10

# Data Generator for Training data
train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(image_size, image_size),
    batch_size=train_batchsize,

```

```

class_mode='categorical')

# Data Generator for Validation data
validation_generator = validation_datagen.flow_from_directory(
validation_dir,
target_size=(image_size, image_size),
batch_size=val_batchsize,
class_mode='categorical',
shuffle=False)

# Compile the model
model.compile(loss='categorical_crossentropy',
optimizer=optimizers.RMSprop(lr=1e-4),
metrics=['acc'])

# Train the Model
history = model.fit_generator(
train_generator,
steps_per_epoch=train_generator.samples/train_generator.batch_size ,
epochs=20,
validation_data=validation_generator,
validation_steps=validation_generator.samples/validation_generator.batch_size,
verbose=1)

# Save the Model
model.save('all_freezed.h5')

# Plot the accuracy and loss curves
acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(len(acc))

plt.plot(epochs, acc, 'b', label='Training acc')
plt.plot(epochs, val_acc, 'r', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()

plt.figure()

plt.plot(epochs, loss, 'b', label='Training loss')
plt.plot(epochs, val_loss, 'r', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()

plt.show()

```

خروجی:

```


Found 600 images belonging to 3 classes.
Found 150 images belonging to 3 classes.
Epoch 1/20
6/6 [=====] - 5s - loss: 3.9083 - acc: 0.4633 -
val_loss: 0.9497 - val_acc: 0.5800
Epoch 2/20
6/6 [=====] - 2s - loss: 1.3948 - acc: 0.5717 -
val_loss: 0.5033 - val_acc: 0.7933

```

```

Epoch 3/20
6/6 [=====] - 2s - loss: 0.2767 - acc: 0.9167 -
val_loss: 0.3672 - val_acc: 0.8733
Epoch 4/20
6/6 [=====] - 2s - loss: 0.4351 - acc: 0.8233 -
val_loss: 1.4450 - val_acc: 0.6667
Epoch 5/20
6/6 [=====] - 2s - loss: 0.3103 - acc: 0.8967 -
val_loss: 0.3865 - val_acc: 0.8667
Epoch 6/20
6/6 [=====] - 2s - loss: 0.2525 - acc: 0.8867 -
val_loss: 0.3298 - val_acc: 0.8733
Epoch 7/20
6/6 [=====] - 2s - loss: 0.0763 - acc: 0.9850 -
val_loss: 0.2922 - val_acc: 0.9000
Epoch 8/20
6/6 [=====] - 2s - loss: 0.3435 - acc: 0.8650 -
val_loss: 0.9127 - val_acc: 0.7200
Epoch 9/20
6/6 [=====] - 2s - loss: 0.1685 - acc: 0.9400 -
val_loss: 0.2946 - val_acc: 0.9000
Epoch 10/20
6/6 [=====] - 2s - loss: 0.0318 - acc: 1.0000 -
val_loss: 0.2864 - val_acc: 0.9067
Epoch 11/20
6/6 [=====] - 2s - loss: 0.0291 - acc: 0.9983 -
val_loss: 0.3485 - val_acc: 0.8733
Epoch 12/20
6/6 [=====] - 2s - loss: 0.0215 - acc: 1.0000 -
val_loss: 0.3003 - val_acc: 0.8933
Epoch 13/20
6/6 [=====] - 2s - loss: 0.0640 - acc: 0.9767 -
val_loss: 3.1314 - val_acc: 0.6133
Epoch 14/20
6/6 [=====] - 2s - loss: 0.6810 - acc: 0.7700 -
val_loss: 0.3014 - val_acc: 0.9000
Epoch 15/20
6/6 [=====] - 2s - loss: 0.0139 - acc: 0.9983 -
val_loss: 0.2988 - val_acc: 0.8933
Epoch 16/20
6/6 [=====] - 2s - loss: 0.0109 - acc: 1.0000 -
val_loss: 0.2988 - val_acc: 0.8933
Epoch 17/20
6/6 [=====] - 2s - loss: 0.0085 - acc: 0.9983 -
val_loss: 0.2928 - val_acc: 0.9000
Epoch 18/20
6/6 [=====] - 2s - loss: 0.0057 - acc: 1.0000 -
val_loss: 0.2994 - val_acc: 0.8933
Epoch 19/20
6/6 [=====] - 2s - loss: 0.0049 - acc: 1.0000 -
val_loss: 0.3011 - val_acc: 0.8933
Epoch 20/20
6/6 [=====] - 2s - loss: 0.0037 - acc: 1.0000 -
val_loss: 0.2972 - val_acc: 0.9000

```



حال برای آزمایش ۴ لایه آخر لایه های convolutional را در نظر می گیریم:

```
from keras.applications import VGG16

#Load the VGG model
vgg_conv = VGG16(weights='imagenet', include_top=False, input_shape=(image_size,
image_size, 3))

# Freeze all the layers
for layer in vgg_conv.layers[:-4]:
    layer.trainable = False

# Check the trainable status of the individual layers
for layer in vgg_conv.layers:
    print(layer, layer.trainable)

from keras import models
from keras import layers
from keras import optimizers

# Create the model
model = models.Sequential()

# Add the vgg convolutional base model
model.add(vgg_conv)

# Add new layers
model.add(layers.Flatten())
model.add(layers.Dense(1024, activation='relu'))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(3, activation='softmax'))

# Show a summary of the model. Check the number of trainable parameters
model.summary()
```

خروجی:

```
<keras.engine.topology.InputLayer object at 0x7f2eec8eeba8> False
<keras.layers.convolutional.Conv2D object at 0x7f2eec8ee320> False
<keras.layers.convolutional.Conv2D object at 0x7f2eec8eed30> False
<keras.layers.pooling.MaxPooling2D object at 0x7f2eeb9551d0> False
<keras.layers.convolutional.Conv2D object at 0x7f2eeb8f9dd8> False
<keras.layers.convolutional.Conv2D object at 0x7f2eec5e3a58> False
<keras.layers.pooling.MaxPooling2D object at 0x7f2eeb906b00> False
<keras.layers.convolutional.Conv2D object at 0x7f2eea3b61d0> False
<keras.layers.convolutional.Conv2D object at 0x7f2eea3b6518> False
<keras.layers.convolutional.Conv2D object at 0x7f2eea3a3b70> False
<keras.layers.pooling.MaxPooling2D object at 0x7f2eea3df630> False
<keras.layers.convolutional.Conv2D object at 0x7f2eea3d6a20> False
<keras.layers.convolutional.Conv2D object at 0x7f2eea3d6828> False
<keras.layers.convolutional.Conv2D object at 0x7f2eea25a908> False
<keras.layers.pooling.MaxPooling2D object at 0x7f2eea27bef0> False
<keras.layers.convolutional.Conv2D object at 0x7f2eea220438> True
<keras.layers.convolutional.Conv2D object at 0x7f2eea220f60> True
<keras.layers.convolutional.Conv2D object at 0x7f2eea1c0cc0> True
<keras.layers.pooling.MaxPooling2D object at 0x7f2eea1e2fd0> True
```



Layer (type)	Output Shape	Param #
vgg16 (Model)	(None, 7, 7, 512)	14714688
flatten_4 (Flatten)	(None, 25088)	0
dense_7 (Dense)	(None, 1024)	25691136
dropout_4 (Dropout)	(None, 1024)	0
dense_8 (Dense)	(None, 3)	3075
Total params: 40,408,899		
Trainable params: 32,773,635		
Non-trainable params: 7,635,264		

در نهایت دوباره مدل را آموزش می‌دهیم و از آن تست می‌گیریم:

```

train_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=20,
    width_shift_range=0.2,
    height_shift_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest')

validation_datagen = ImageDataGenerator(rescale=1./255)

# Change the batchsize according to your system RAM
train_batchsize = 50
val_batchsize = 10

# Data Generator for Training data
train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(image_size, image_size),
    batch_size=train_batchsize,
    class_mode='categorical')

# Data Generator for Validation data
validation_generator = validation_datagen.flow_from_directory(
    validation_dir,
    target_size=(image_size, image_size),
    batch_size=val_batchsize,
    class_mode='categorical',
    shuffle=False)

# Compile the model

```

```

model.compile(loss='categorical_crossentropy',
optimizer=optimizers.RMSprop(lr=1e-4),
metrics=['acc'])

# Train the Model
# NOTE that we have multiplied the steps_per_epoch by 2. This is because we
are using data augmentation.
history = model.fit_generator(
train_generator,
steps_per_epoch=2*train_generator.samples/train_generator.batch_size ,
epochs=40,
validation_data=validation_generator,
validation_steps=validation_generator.samples/validation_generator.batch_size,
verbose=1)

# Save the Model
model.save('da_last4_layers.h5')

# Plot the accuracy and loss curves
acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(len(acc))

plt.plot(epochs, acc, 'b', label='Training acc')
plt.plot(epochs, val_acc, 'r', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()

plt.figure()

plt.plot(epochs, loss, 'b', label='Training loss')
plt.plot(epochs, val_loss, 'r', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()

plt.show()

```

خروجی:

```

Found 600 images belonging to 3 classes.
Found 150 images belonging to 3 classes.
Epoch 1/40
24/24 [=====] - 8s - loss: 1.0960 - acc: 0.5783 -
val_loss: 0.3230 - val_acc: 0.9200
Epoch 2/40
24/24 [=====] - 7s - loss: 0.4048 - acc: 0.8558 -
val_loss: 0.1907 - val_acc: 0.9533
Epoch 3/40
24/24 [=====] - 8s - loss: 0.1802 - acc: 0.9392 -
val_loss: 0.2353 - val_acc: 0.9333
Epoch 4/40
24/24 [=====] - 7s - loss: 0.1277 - acc: 0.9517 -
val_loss: 0.1922 - val_acc: 0.9667
Epoch 5/40
24/24 [=====] - 8s - loss: 0.1963 - acc: 0.9417 -
val_loss: 0.1390 - val_acc: 0.9800

```

```



Epoch 6/40
24/24 [=====] - 7s - loss: 0.0414 - acc: 0.9908 -
val_loss: 0.3272 - val_acc: 0.9533
Epoch 7/40
24/24 [=====] - 7s - loss: 0.1018 - acc: 0.9717 -
val_loss: 0.1854 - val_acc: 0.9733
Epoch 8/40
24/24 [=====] - 7s - loss: 0.0733 - acc: 0.9792 -
val_loss: 0.1771 - val_acc: 0.9667
Epoch 9/40
24/24 [=====] - 7s - loss: 0.0478 - acc: 0.9892 -
val_loss: 0.3292 - val_acc: 0.9133
Epoch 10/40
24/24 [=====] - 7s - loss: 0.0138 - acc: 0.9958 -
val_loss: 0.2593 - val_acc: 0.9600
Epoch 11/40
24/24 [=====] - 7s - loss: 0.0788 - acc: 0.9800 -
val_loss: 0.2195 - val_acc: 0.9800
Epoch 12/40
24/24 [=====] - 7s - loss: 0.0329 - acc: 0.9925 -
val_loss: 0.2266 - val_acc: 0.9733
Epoch 13/40
24/24 [=====] - 7s - loss: 2.2249e-04 - acc: 1.0000
- val_loss: 0.2333 - val_acc: 0.9733
Epoch 14/40
24/24 [=====] - 7s - loss: 0.1973 - acc: 0.9792 -
val_loss: 0.1446 - val_acc: 0.9800
Epoch 15/40
24/24 [=====] - 7s - loss: 0.0024 - acc: 1.0000 -
val_loss: 0.1517 - val_acc: 0.9733
Epoch 16/40
24/24 [=====] - 7s - loss: 2.6369e-04 - acc: 1.0000
- val_loss: 0.1819 - val_acc: 0.9800
Epoch 17/40
24/24 [=====] - 8s - loss: 0.2217 - acc: 0.9742 -
val_loss: 0.1965 - val_acc: 0.9667
Epoch 18/40
24/24 [=====] - 7s - loss: 0.0028 - acc: 1.0000 -
val_loss: 0.2090 - val_acc: 0.9800
Epoch 19/40
24/24 [=====] - 8s - loss: 6.4382e-04 - acc: 1.0000
- val_loss: 0.1917 - val_acc: 0.9800
Epoch 20/40
24/24 [=====] - 7s - loss: 0.1931 - acc: 0.9758 -
val_loss: 0.1457 - val_acc: 0.9667
Epoch 21/40
24/24 [=====] - 7s - loss: 4.4634e-04 - acc: 1.0000
- val_loss: 0.1774 - val_acc: 0.9667
Epoch 22/40
24/24 [=====] - 7s - loss: 1.3386e-04 - acc: 1.0000
- val_loss: 0.2766 - val_acc: 0.9733
Epoch 23/40
24/24 [=====] - 7s - loss: 0.0733 - acc: 0.9900 -
val_loss: 0.1287 - val_acc: 0.9800
Epoch 24/40
24/24 [=====] - 7s - loss: 2.8142e-04 - acc: 1.0000
- val_loss: 0.1496 - val_acc: 0.9867

```

```

Epoch 25/40
24/24 [=====] - 7s - loss: 4.9226e-05 - acc: 1.0000
- val_loss: 0.1093 - val_acc: 0.9933
Epoch 26/40
24/24 [=====] - 7s - loss: 1.9475e-06 - acc: 1.0000
- val_loss: 0.1091 - val_acc: 0.9933
Epoch 27/40
24/24 [=====] - 7s - loss: 5.5070e-06 - acc: 1.0000
- val_loss: 0.1217 - val_acc: 0.9800
Epoch 28/40
24/24 [=====] - 7s - loss: 0.2148 - acc: 0.9775 -
val_loss: 0.1222 - val_acc: 0.9867
Epoch 29/40
24/24 [=====] - 7s - loss: 0.0671 - acc: 0.9883 -
val_loss: 0.2639 - val_acc: 0.9600
Epoch 30/40
24/24 [=====] - 7s - loss: 0.0088 - acc: 0.9975 -
val_loss: 0.1864 - val_acc: 0.9733
Epoch 31/40
24/24 [=====] - 8s - loss: 0.0186 - acc: 0.9975 -
val_loss: 0.1440 - val_acc: 0.9800
Epoch 32/40
24/24 [=====] - 7s - loss: 0.1914 - acc: 0.9775 -
val_loss: 0.1286 - val_acc: 0.9800
Epoch 33/40
24/24 [=====] - 8s - loss: 0.0027 - acc: 0.9983 -
val_loss: 0.1252 - val_acc: 0.9867
Epoch 34/40
24/24 [=====] - 7s - loss: 0.1001 - acc: 0.9892 -
val_loss: 0.1778 - val_acc: 0.9667
Epoch 35/40
24/24 [=====] - 7s - loss: 0.0014 - acc: 1.0000 -
val_loss: 0.2393 - val_acc: 0.9733
Epoch 36/40
24/24 [=====] - 7s - loss: 4.4375e-05 - acc: 1.0000
- val_loss: 0.1761 - val_acc: 0.9867
Epoch 37/40
24/24 [=====] - 7s - loss: 1.3121e-06 - acc: 1.0000
- val_loss: 0.1838 - val_acc: 0.9867
Epoch 38/40
24/24 [=====] - 7s - loss: 7.2186e-07 - acc: 1.0000
- val_loss: 0.1935 - val_acc: 0.9867
Epoch 39/40
24/24 [=====] - 7s - loss: 8.2086e-07 - acc: 1.0000
- val_loss: 0.2051 - val_acc: 0.9800
Epoch 40/40
24/24 [=====] - 7s - loss: 0.3185 - acc: 0.9717 -
val_loss: 0.1670 - val_acc: 0.9800

```



حال با تست می‌کنیم که عکس‌ها را درست پیش‌بینی می‌کند یا نه:

```
# Create a generator for prediction
validation_generator = validation_datagen.flow_from_directory(
    validation_dir,
    target_size=(image_size, image_size),
    batch_size=val_batchsize,
    class_mode='categorical',
    shuffle=False)

# Get the filenames from the generator
fnames = validation_generator.filenames

# Get the ground truth from generator
ground_truth = validation_generator.classes

# Get the label to class mapping from the generator
label2index = validation_generator.class_indices

# Getting the mapping from class index to class label
idx2label = dict((v,k) for k,v in label2index.items())

# Get the predictions from the model using the generator
predictions = model.predict_generator(validation_generator,
steps=validation_generator.samples/validation_generator.batch_size,verbose=1)
predicted_classes = np.argmax(predictions,axis=1)

errors = np.where(predicted_classes != ground_truth)[0]
print("No of errors = {}".format(len(errors),validation_generator.samples))

# Show the errors
for i in range(len(errors)):
    pred_class = np.argmax(predictions[errors[i]])
    pred_label = idx2label[pred_class]

    title = 'Original label:{}, Prediction :{}, confidence : {:.3f}'.format(
        fnames[errors[i]].split('/')[0],
        pred_label,
        predictions[errors[i]][pred_class])

    original = load_img('{}{}'.format(validation_dir,fnames[errors[i]]))
    plt.figure(figsize=[7,7])
    plt.axis('off')
    plt.title(title)
    plt.imshow(original)
    plt.show()
```

خروجی:

```
Found 150 images belonging to 3 classes.
15/15 [=====] - 0s
No of errors = 3/150
```

Original label:pumpkin, Prediction :watermelon, confidence : 0.999



Original label:tomato, Prediction :watermelon, confidence : 1.000




Original label:watermelon, Prediction :tomato, confidence : 0.432



## سوال ۵

برای این سوال یک ارائه به صورت pdf تهیه کردم و آنرا مطالعه کردم و از نتایج آن فهمیدم بر روی عکس ها پردازش می کند تا feature هایی که میخواهد را بدست بیاورد که نتایج آن مثل شکل های زیر است:



## سوال ۶

در این سوال کافی است دستورات زیر اجرا شود:

```
print('Checking the Training on a Single Batch...')
with tf.Session() as session:
    # Initializing the variables
    session.run(tf.global_variables_initializer())

    # Training cycle
    for epoch in range(epochs):
        batch_i = 1
        for batch_features, batch_labels in \
            load_preprocess_training_batch(batch_i, batch_size):
            session.run(optimizer, feed_dict=\
                {x:batch_features, y:batch_labels, keep_prob:keep_probability})
            print('Epoch {:>2}, CIFAR-10 Batch {}: '.format(epoch + 1, batch_i), end='')
            print_stats(session, batch_features, batch_labels, cost, accuracy)
```

خروجی آن نیز به شکل زیر است:

```
Checking the Training on a Single Batch...
Epoch  1, CIFAR-10 Batch 1: batch loss is :  2.04934
batch_accuracy accuracy is : 0.3524
```



Epoch 2, CIFAR-10 Batch 1: batch loss is : 1.80641  
batch\_accuracy accuracy is : 0.4126  
Epoch 3, CIFAR-10 Batch 1: batch loss is : 1.54692  
batch\_accuracy accuracy is : 0.4486  
Epoch 4, CIFAR-10 Batch 1: batch loss is : 1.27744  
batch\_accuracy accuracy is : 0.4792  
Epoch 5, CIFAR-10 Batch 1: batch loss is : 1.13442  
batch\_accuracy accuracy is : 0.4928  
Epoch 6, CIFAR-10 Batch 1: batch loss is : 0.92034  
batch\_accuracy accuracy is : 0.496  
Epoch 7, CIFAR-10 Batch 1: batch loss is : 0.806121  
batch\_accuracy accuracy is : 0.4896  
Epoch 8, CIFAR-10 Batch 1: batch loss is : 0.745746  
batch\_accuracy accuracy is : 0.5048  
Epoch 9, CIFAR-10 Batch 1: batch loss is : 0.615836  
batch\_accuracy accuracy is : 0.531  
Epoch 10, CIFAR-10 Batch 1: batch loss is : 0.497367  
batch\_accuracy accuracy is : 0.5288  
Epoch 11, CIFAR-10 Batch 1: batch loss is : 0.442382  
batch\_accuracy accuracy is : 0.531  
Epoch 12, CIFAR-10 Batch 1: batch loss is : 0.381339  
batch\_accuracy accuracy is : 0.55  
Epoch 13, CIFAR-10 Batch 1: batch loss is : 0.365281  
batch\_accuracy accuracy is : 0.5634  
Epoch 14, CIFAR-10 Batch 1: batch loss is : 0.307827  
batch\_accuracy accuracy is : 0.5432  
Epoch 15, CIFAR-10 Batch 1: batch loss is : 0.26708  
batch\_accuracy accuracy is : 0.571  
Epoch 16, CIFAR-10 Batch 1: batch loss is : 0.222317  
batch\_accuracy accuracy is : 0.577  
Epoch 17, CIFAR-10 Batch 1: batch loss is : 0.191946  
batch\_accuracy accuracy is : 0.5728  
Epoch 18, CIFAR-10 Batch 1: batch loss is : 0.142695  
batch\_accuracy accuracy is : 0.5794  
Epoch 19, CIFAR-10 Batch 1: batch loss is : 0.176674  
batch\_accuracy accuracy is : 0.5508  
Epoch 20, CIFAR-10 Batch 1: batch loss is : 0.163542  
batch\_accuracy accuracy is : 0.579  
Epoch 21, CIFAR-10 Batch 1: batch loss is : 0.171183  
batch\_accuracy accuracy is : 0.5622  
Epoch 22, CIFAR-10 Batch 1: batch loss is : 0.132356  
batch\_accuracy accuracy is : 0.5728  
Epoch 23, CIFAR-10 Batch 1: batch loss is : 0.10712  
batch\_accuracy accuracy is : 0.5704  
Epoch 24, CIFAR-10 Batch 1: batch loss is : 0.101802  
batch\_accuracy accuracy is : 0.5618  
Epoch 25, CIFAR-10 Batch 1: batch loss is : 0.108635  
batch\_accuracy accuracy is : 0.5678  
Epoch 26, CIFAR-10 Batch 1: batch loss is : 0.112882  
batch\_accuracy accuracy is : 0.5596  
Epoch 27, CIFAR-10 Batch 1: batch loss is : 0.0682271  
batch\_accuracy accuracy is : 0.5884  
Epoch 28, CIFAR-10 Batch 1: batch loss is : 0.0640762  
batch\_accuracy accuracy is : 0.5856  
Epoch 29, CIFAR-10 Batch 1: batch loss is : 0.0570013  
batch\_accuracy accuracy is : 0.5866  
Epoch 30, CIFAR-10 Batch 1: batch loss is : 0.0763902

```

batch_accuracy accuracy is : 0.583
Epoch 31, CIFAR-10 Batch 1: batch loss is : 0.0564564
batch_accuracy accuracy is : 0.571
Epoch 32, CIFAR-10 Batch 1: batch loss is : 0.0532774
batch_accuracy accuracy is : 0.588
Epoch 33, CIFAR-10 Batch 1: batch loss is : 0.0437081
batch_accuracy accuracy is : 0.5838
Epoch 34, CIFAR-10 Batch 1: batch loss is : 0.0328673
batch_accuracy accuracy is : 0.5664
Epoch 35, CIFAR-10 Batch 1: batch loss is : 0.0359792
batch_accuracy accuracy is : 0.5884
Epoch 36, CIFAR-10 Batch 1: batch loss is : 0.0376761
batch_accuracy accuracy is : 0.5898
Epoch 37, CIFAR-10 Batch 1: batch loss is : 0.0326239
batch_accuracy accuracy is : 0.5694
Epoch 38, CIFAR-10 Batch 1: batch loss is : 0.0366496
batch_accuracy accuracy is : 0.589
Epoch 39, CIFAR-10 Batch 1: batch loss is : 0.0224001
batch_accuracy accuracy is : 0.582
Epoch 40, CIFAR-10 Batch 1: batch loss is : 0.0224597
batch_accuracy accuracy is : 0.5896
Epoch 41, CIFAR-10 Batch 1: batch loss is : 0.0199231
batch_accuracy accuracy is : 0.5834
Epoch 42, CIFAR-10 Batch 1: batch loss is : 0.0231632
batch_accuracy accuracy is : 0.5878
Epoch 43, CIFAR-10 Batch 1: batch loss is : 0.0191447
batch_accuracy accuracy is : 0.5932
Epoch 44, CIFAR-10 Batch 1: batch loss is : 0.0206479
batch_accuracy accuracy is : 0.5776
Epoch 45, CIFAR-10 Batch 1: batch loss is : 0.0263408
batch_accuracy accuracy is : 0.5718
Epoch 46, CIFAR-10 Batch 1: batch loss is : 0.0125673
batch_accuracy accuracy is : 0.5802
Epoch 47, CIFAR-10 Batch 1: batch loss is : 0.0197601
batch_accuracy accuracy is : 0.5828
Epoch 48, CIFAR-10 Batch 1: batch loss is : 0.0156334
batch_accuracy accuracy is : 0.5908
Epoch 49, CIFAR-10 Batch 1: batch loss is : 0.011323
batch_accuracy accuracy is : 0.5846
Epoch 50, CIFAR-10 Batch 1: batch loss is : 0.0132246
batch_accuracy accuracy is : 0.5784

```

حال مدل fully trained آن را بدست می آورم:

```

save_model_path = './image_classification'

print('Training...')
with tf.Session() as session:
    # Initializing the variables
    session.run(tf.global_variables_initializer())

    # Training cycle
    for epoch in range(epochs):
        # Loop over all batches
        n_batches = 5

```

```

for batch_i in range(1, n_batches + 1):
    for batch_features, batch_labels in \
        load_preprocess_training_batch(batch_i, batch_size):
        session.run(optimizer, feed_dict=\
            {x:batch_features, y:batch_labels, keep_prob:keep_probability})

# Save Model
saver = tf.train.Saver()
save_path = saver.save(session, save_model_path)

```

حال نوبت به آن رسید که تست کنیم تصاویری که می‌دهیم درست تشخیص می‌دهد یا نه.

```

%matplotlib inline
%config InlineBackend.figure_format = 'retina'

import tensorflow as tf
import pickle
import random
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelBinarizer

def display_image_predictions(features, labels, predictions):
    label_binarizer = LabelBinarizer()
    label_binarizer.fit(range(LABELS_COUNT))
    label_ids = label_binarizer.inverse_transform(np.array(labels))

    fig, axes = plt.subplots(nrows=4, ncols=2)
    fig.tight_layout()
    fig.suptitle('Softmax Predictions', fontsize=20, y=1.1)

    n_predictions = 3
    margin = 0.05
    ind = np.arange(n_predictions)
    width = (1. - 2. * margin) / n_predictions

    for image_i, (feature, label_id, pred_indicies, pred_values) \
        in enumerate(zip(features, label_ids, predictions.indices,
            predictions.values)):
        pred_names = [LABEL_NAMES[pred_i] for pred_i in pred_indicies]
        correct_name = LABEL_NAMES[label_id]

        axes[image_i][0].imshow(feature)
        axes[image_i][0].set_title(correct_name)
        axes[image_i][0].set_axis_off()

        axes[image_i][1].barh(ind + margin, pred_values[:-1], width)
        axes[image_i][1].set_yticks(ind + margin)
        axes[image_i][1].set_yticklabels(pred_names[:-1])
        axes[image_i][1].set_xticks([0, 0.5, 1.0])

    save_model_path = './image_classification'
    n_samples = 4
    top_n_predictions = 3

def test_model():

```

```

"""
Test the saved model against the test dataset
"""

test_features, test_labels = pickle.load(open('preprocess_training.p',
mode='rb'))
loaded_graph = tf.Graph()

with tf.Session(graph=loaded_graph) as sess:
    # Load model
    loader = tf.train.import_meta_graph(save_model_path + '.meta')
    loader.restore(sess, save_model_path)

    # Get Tensors from loaded model
    loaded_x = loaded_graph.get_tensor_by_name('x:0')
    loaded_y = loaded_graph.get_tensor_by_name('y:0')
    loaded_keep_prob = loaded_graph.get_tensor_by_name('keep_prob:0')
    loaded_logits = loaded_graph.get_tensor_by_name('logits:0')
    loaded_acc = loaded_graph.get_tensor_by_name('accuracy:0')

    # Get accuracy in batches for memory limitations
    test_batch_acc_total = 0
    test_batch_count = 0

    for train_feature_batch, train_label_batch in \
        batch_features_labels(test_features, test_labels, batch_size):
        test_batch_acc_total += sess.run(
            loaded_acc,
            feed_dict={loaded_x: train_feature_batch, \
                loaded_y: train_label_batch, loaded_keep_prob: 1.0})
        test_batch_count += 1

    print('Testing Accuracy: {}\n'.format(test_batch_acc_total/test_batch_count))

    # Print Random Samples
    random_test_features, random_test_labels = \
        tuple(zip(*random.sample(list(zip(test_features, test_labels)), n_samples)))

    random_test_predictions = sess.run(
        tf.nn.top_k(tf.nn.softmax(loaded_logits), top_n_predictions),
        feed_dict={loaded_x: random_test_features, \
            loaded_y: random_test_labels, loaded_keep_prob: 1.0})
    display_image_predictions(random_test_features, \
        random_test_labels, random_test_predictions)

test_model()


```

نتیجه ای که بدست آمد به شرح زیر است:


Testing Accuracy: 0.6734375

## Softmax Predictions


bird





ship



airplane



deer



به پایان آمد این دفتر، حکایت همچنان باقی

کتاب بلغ منی حیا معرضاً عنی

نکویم نسبتی دارم به نزدیکان درگاهت

اخلاقی و احبابی ذوا من جدایی

به صد دفتر شاید گفت شرح الحال مشتاقی

ان افضل ماتری آنی علی عهدی و میثاقی

که خود را بر تو می بندم به سالوسی و زرقی

مریض العشق لایسری ولای شکوای الراتی

تو را گر خواب می گیرد نه صاحب درد عشاقی	نشان عاشق آن باشد که شب باروز پیوندد
اما انت الذی تتقی فحین السم تریاقی	قم املا واسقنی کلسا ووع مافیہ مسموما
مرا بگذارتا حیران بماند چشم در ساقی	قدح چون دور ما باشد به هشیار ان مجلس ده
انا الجحون لا اعبا باحراق و اغراق	سعی فی حکمی الشانی و لما یدر ماشانی
مگر نفس ملک باشد بدین پاکیزه اخلاقی	مگر شمس فلک باشد بدین فرخنده دیداری
وهذا الطبی فی شیراز یسینی باحداق	لقتیت الاسد فی الغابات لا تقوی علی صیدی
بمیرد تشنه مستقی و دریا بچنان باقی	نه حسنت آخری دارد، نه سعدی را سخن پایان

پایان